DANIEL DOBROVOLSKIY

Austin, TX | 346-652-4563 | danieldob200@gmail.com | linkedin.com/in/danieldobrovolskiy | dramarama.vercel.app

SUMMARY

Full-stack software engineer shipping production code since 2021, with secure software for a high-compliance environment and, since 2025, production AI systems: RAG and semantic search pipelines, agentic workflows with validation controls, and LLM guardrails. Team AI Champion guiding engineers on safe, effective AI-assisted development, with startup experience at Neurohive and DramaRama.

TECHNICAL SKILLS

GenAI & Retrieval: RAG, embeddings, semantic and vector search (OpenSearch), indexing and ranking, LangChain, LangGraph, LLM tool calling, structured outputs, streaming, prompt and context engineering, guardrails, grounding
Languages & Web: Python, Java (Spring Boot), TypeScript, JavaScript, SQL, C#; React, Next.js, FastAPI, REST, WebSockets, PostgreSQL
Cloud, DevOps & Security: AWS (Lambda, S3, EventBridge, DynamoDB, EC2, SAM IaC), Docker, CI/CD, serverless, microservices; RBAC, tenant isolation, JWT/OAuth, pen-test remediation
AI-Assisted Development: modular agentic coding workflows in Claude Code, Codex, and Devin with shared context, templates, and test-based validation

EXPERIENCE

Technical Support Consultant (Full-Stack Software Engineer) | Baker Tilly | Austin, TX (Remote) **Sep 2024 – Present**
Startup-style product team working directly with a firm partner; ships rapid applications prioritized by usage metrics.

- Appointed team AI Champion (2025): teach engineers AI-assisted workflows, run AI Lunch & Learns, and represent the team with the firm's AI enablement group on tool selection.
- Lead developer of the Domestic Content web app; build multi-tenant compliance software for Inflation Reduction Act tax-credit programs in Python and PostgreSQL: S3-to-PostgreSQL ingestion, a penalty-calculating rules engine, and Tableau views, with tenant isolation, RBAC, and pen-test remediation.
- Building a documentation-grounded AI chatbot that answers portal users' how-to questions, shaped through user interviews and the platform's lead PM, who supplies the knowledge base.
- Built an internal workbench to edit and trace each chatbot step (prompts, code, retrieved context) and compare LLMs.
- Designed a chain of modular AI workflows that take a ticket from requirements to QA: an interview saves business context to a ticket workspace, a technical pass builds on it, and unclear requirements go to the PM or a senior engineer before coding starts.
- The workflows then write unit or integration tests and mocks from team templates before the code, so AI-generated code is built to meet the agreed behavior, and draft a QA test plan, creating any missing test data; a shared script library keeps token use low.

AI Engineer (Contract) | Neurohive (AI startup) | Remote **Jul 2025 – Jan 2026**

- Built agentic RAG workflows (LangGraph, LangChain) for a global oil and gas company, answering questions over large structured and unstructured datasets on a small, fast-moving team.
- With the architect, built a serverless multi-agent pipeline: a router Lambda sizes each question, a context manager fans out specialist Lambdas (dataset, web, quantitative search) to S3, and EventBridge triggers a finalizer that streams the answer over WebSocket.
- Built retrieval with embedding-based semantic search on OpenSearch, dataset-linked tools ranked to pick the most relevant data for each question, and token-efficiency techniques; reviewed junior developers' code.

Software Engineer (Contract via Pinnacle Group) | JPMorgan Chase & Co. | Houston, TX **Oct 2022 – Jan 2023**

- Migrated Java/Spring Boot microservices to AWS with Docker and CI/CD; built React UIs with design, product, and BA teams.

Freelance Software Developer | Self-employed | danban451.github.io/portfolio **2021 – 2024**

- Hand-coded and launched websites for small-business clients (shown in my portfolio), working with a UX designer; set up Hygraph CMS so clients could manage their own content.
- Learned the fundamentals by hand: vanilla JavaScript and direct DOM manipulation, then React with reusable components, Redux and MobX state management, and load-time optimizations; also Angular and MVC patterns.

PROJECTS

DramaRama | Founder & AI Engineer, Whimsy & Rigor LLC (startup) | dramarama.vercel.app **2025 – Present**

- Built a multi-agent LLM platform (FastAPI, Claude, Next.js, TypeScript, PostgreSQL, AWS): 12 advisor agents plus an orchestrator, LLM-judged guardrails, and per-call traces with replayable sessions.
- Validate demand with unbiased in-person user interviews in Austin and feedback from the framework's author, a C-suite PR advisor, and a founder coach; six major pivots driven by testing.

TRAINING & CERTIFICATIONS

Formal training in software engineering concepts: 37 professional certifications (LinkedIn Learning, Coursera, Udemy).